

DataGovOps DES 2026

Speaker: Piyush Ahuja, Director, Tata Digital

Session title: *Governing GenAI: DataGovOps When AI Is Your Biggest Data Consumer*

Format: 15-minute standalone talk · 5 min QnA · 30s setup · 30s exit

Theme alignment: Data Foundations · DataOps · Responsible AI

Audience: C-suite leaders, data and AI leaders, technology executives, architects, senior engineers, and governance/risk stakeholders

Organizer Summary

This session is designed as a high-energy, practitioner-led talk for senior data, AI, technology, and business audiences. It introduces a clear category shift: enterprise data governance was built for human data consumers, but GenAI creates **agentic data consumers** that operate continuously, non-deterministically, and at machine speed.

The talk is vendor-neutral, non-promotional, and uses synthetic enterprise scenarios rather than company-specific examples. It introduces **DataGovOps** as a memorable operating framework for governing AI data consumption through policy-as-code, intent-based access, consent-as-runtime-control, source-to-output lineage, and governed AI-assisted classification.

The session is structured to give leaders a strategic frame, architects an implementation model, and practitioners three concrete actions they can take the following week.

Purpose and Category Shift

For the last two decades, most enterprise governance programs assumed the primary data consumer was human: an analyst writing SQL, a business user opening a dashboard, a data scientist requesting access, or an engineer building a pipeline.

GenAI changes that assumption.

AI agents consume data through prompts, embeddings, vector stores, tools, memory, API calls, and downstream outputs. They do this continuously, at machine speed, and often without the deterministic behavior that traditional controls were designed around.

Category shift: Governance was built for human data consumers. DataGovOps is governance for agentic data consumers.

Core thesis:

Your biggest data consumer is no longer human. Your governance was built for humans. That gap is your next incident.

What the Audience Will Learn

Attendees will leave with a practical framework for governing GenAI systems when the primary data consumer is no longer a human analyst, but an AI agent operating continuously and at machine speed.

They will learn:

1. Why traditional data governance breaks under agentic workloads

Governance models designed for human, deterministic, slow-moving consumers struggle when AI systems consume prompts, embeddings, memory, tool outputs, and source data in real time.

2. The DataGovOps operating model

A four-part framework for AI-era governance: policy as code, least privilege per intent, source-to-output lineage, and governed AI-assisted classification.

3. How to translate governance into engineering controls

The talk shows how consent, access, classification, lineage, and auditability can move from documents and approvals into runtime systems, CI/CD checks, and traceable data paths.

4. Three actions to take after the summit

Audit one AI pipeline end-to-end, move one governance policy into code, and map access by intent rather than identity alone.

Audience Fit

The session is designed for a mixed seniority audience. Each segment works at both strategic and implementation levels.

Audience segment	Primary value	What they will engage with most
C-suite / VPs	Strategic framing: governance as an enabler of responsible AI velocity, not a brake.	The category shift to agentic data consumers, the Governance 1.0 failure modes, and the regulatory readiness framing.
Directors / Heads of Data and AI	Operating model: how governance, data platforms, and AI delivery need to converge.	The DataGovOps framework, intent-based access, and source-to-output lineage.
Architects / Senior Engineers	Implementation patterns that can be prototyped or folded into existing platforms.	Policy-as-code, tag propagation through embeddings, lineage traces, and eval gates for governance AI.
Compliance / Risk leaders	Clear translation between technical architecture and governance evidence.	Consent as runtime control, incident traceability, and synthetic proof walkthrough.

Audience Experience Design

The session is designed as a high-energy, practitioner-led talk rather than a policy lecture. It uses concrete enterprise scenarios, minimal text-heavy slides, and a clear repeated thesis to help the audience remember the framework.

The delivery follows four principles:

- **Scenario first, label second** — every major idea starts with a recognizable AI governance failure before introducing the DataGovOps pattern.
 - **Memorable vocabulary** — phrases such as “Governance 1.0,” “least privilege per intent,” and “policy in Git, not in Confluence” are used to make the framework portable.
 - **Proof over theory** — a short, synthetic pre-recorded walkthrough demonstrates how source-to-output lineage can reduce AI incident tracing from weeks to minutes.
 - **Actionable close** — the session ends with three practical steps attendees can apply without needing a major platform transformation.
-

Session Logistics

Item	Detail
Format	20 minute standalone talk, including transition time
Content duration	15 minutes
Q&A	Will support 3–5 minutes if needed
Audience	C-suite leaders, VPs, Directors, Heads of Data/AI/Technology, architects, and senior engineers
Theme	Data Foundations, DataOps, Data Governance, GenAI Readiness
Slide style	Minimal, visual-led, low text density